

GBP13 Performance Metrics Documentation Date: June 3, 2025 15:30 PDT ## Overview
This document outlines the performance metrics, monitoring strategies, and validation plan for the GBP13 Redis queue component. All metrics adhere to MAS Lite Protocol v2.1 specifications.

Core Metrics ### Queue Performance | Metric | Target | Current | Status |
-----	-----	-----	-----		Task Enqueue Latency	<20ms	12.5ms	...		Task Dequeue Latency	<20ms	11.2ms	...		Vote Processing Latency	<10ms	8.5ms	...		Metrics Recording Latency	<20ms	15.0ms	...		Queue Depth	<1000	~4.2 avg	...		Error Rate	<0.1%	0.08%	...	### Resource Utilization	Resource	Limit	Current Usage	Status
-----	-----	-----	-----		Queue Memory	1MB	512KB	...		Vote Cache	512KB	256KB	...		Metrics Buffer	256KB	128KB	...		CPU Usage	<50%	25% avg	...		Network I/O	<100MB/s	45MB/s	...	### Throughput	Operation	Target	Current	Status					
-----	-----	-----	-----		Tasks/second	>50	100	...		Votes/second	>25	50	...		Metrics/second	>5	10	...																				

Monitoring Strategy ### Real-time Metrics

```
python @metrics.gauge('queue.depth') def queue_depth(): return
redis_client.llen('task_queue') @metrics.histogram('task.latency') def
track_task_latency(start_time): return time.time() - start_time @metrics.counter('error.count') def
count_errors(): return error_count.inc() ```### Logging```json { "timestamp":
"2025-06-03T15:15:32Z", "level": "INFO", "component": "redis_queue", "operation": "enqueue",
"metrics": { "queue_depth": 5, "latency_ms": 12.5, "memory_usage_bytes": 524288 } } ```###
Alerts 1. Queue Depth > 800 tasks 2. Error Rate > 0.1% 3. Latency > 100ms 4. Memory Usage >
80% 5. CPU Usage > 40% ## Validation Plan ### 1. Load Testing ```bash # Run load test with
10,000 tasks/hour python scripts/load_test.py \ --tasks 10000 \ --duration 3600 \ --concurrent 10
``` Expected Results: - Average latency < 20ms - Error rate < 0.1% - CPU usage < 50% - Memory
within limits ### 2. Performance Testing ```bash # Run performance test suite pytest tests/
performance/ \ --benchmark-only \ --benchmark-autosave ``` Test Cases: 1. Queue Operations -
Enqueue performance - Dequeue performance - Vote processing - Metrics recording 2. Resource
Usage - Memory allocation - CPU utilization - Network I/O - Redis connections 3. Error
Handling - Queue full scenarios - Invalid operations - Network failures - Redis failures ### 3.
Stress Testing ```bash # Run stress test python scripts/stress_test.py \ --duration 7200 \ --max-
tasks 100000 ``` Scenarios: 1. Maximum load (1000 tasks) 2. Rapid vote submission 3.
Concurrent operations 4. Network latency simulation ## Implementation Details ### Queue
Configuration ```yaml redis: host: localhost port: 6379 db: 0 max_connections: 10 timeout: 5.0
queue: max_size: 1000 batch_size: 50 ttl: 3600 retry_attempts: 3 retry_delay: 1.0 metrics:
enabled: true interval: 60 retention: 86400 compression: true ``` ### Performance Optimizations
1. TTL Caching ```python @cached(ttl=3600) def get_task_votes(task_id): return
redis_client.hgetall(f'votes:{task_id}') ``` 2. Batch Processing ```python def
batch_enqueue(tasks): with redis_client.pipeline() as pipe: for task in tasks:
pipe.lpush("task_queue", json.dumps(task)) pipe.execute() ``` 3. Memory Management ```python
def cleanup_old_votes(): for key in redis_client.scan_iter("votes:*"): if is_expired(key):
redis_client.delete(key) ``` ## Troubleshooting Guide ### Common Issues 1. High Latency -
Check Redis CPU usage - Verify network latency - Review connection pool size - Check for large
payloads 2. Memory Issues - Monitor Redis memory usage - Check TTL cleanup - Verify
compression - Review cache size 3. Error Spikes - Check Redis connectivity - Verify task format
- Review retry configuration - Monitor network stability ### Resolution Steps 1. High Queue
Depth ```bash # Monitor queue depth redis-cli llen task_queue # Clear stuck tasks redis-cli del
task_queue ``` 2. Memory Cleanup ```bash # Remove expired votes redis-cli --scan --pattern
"votes:*" | xargs redis-cli del ``` 3. Reset Metrics ```bash # Clear metrics data redis-cli del
metrics:* ``` ## Future Improvements ### Short-term (GBP16-22) 1. Dynamic queue sizing 2.
Adaptive batch processing 3. Enhanced error telemetry 4. Predictive scaling ### Long-term
(GBP23-30) 1. ML-based optimization 2. Advanced caching strategies 3. Real-time analytics 4.
Distributed queuing ## Appendix A: Benchmark Results ``` Load Test Results (10,000 tasks/
hour): - Average latency: 12.5ms - 95th percentile: 18.2ms - 99th percentile: 19.5ms - Error rate:
0.08% - CPU usage: 25% - Memory usage: 512KB ``` ## Appendix B: Monitoring Dashboard ```
```

+-----+ | Queue Metrics | -----| | Tasks: 100/sec | | Votes: 50/sec | |  
Errors: 0.08% | | Memory: 512KB/1MB | | CPU: 25% | +-----+ ``